Verified Search with Neural Ranking and Executable Claims

Rules decide validity; holographic memory stores what is learned; neural models rank what remains; claims stay bound to evidence

Anonymous authors

Abstract. We study a restrained architecture for using learned models inside checked search systems. Three components compose: a verifier-grounded framework that separates validity from preference through certificate-checked symbolic rules, a holographic memory substrate that stores rules as fixed-dimension vector bindings with 1/sqrt(N) retrieval fidelity through a universal continuous primitive. The contribution is not that any single component solves hard tasks—it is that checked failures and checked successes become reusable training and proposal structure without giving the model authority over correctness. The clean anchor is arithmetic sequence filtering: in the audited next-prime instance, a learned modular wheel achieves 100% top-1 accuracy with zero false rejections on 30000 evaluation rows (the verified learned filter raises top-1 ranking to 68.87% from 45.12% for the hand-coded wheel and 9.13% raw, with 0 verifier-counted true-answer deletions.). Holographic storage encodes these certified rules as vector bindings with fixed O(d) memory

holographic ranker achieves 48.5% accuracy with zero trainable parameters versus 16.0% Every empirical claim in this paper is mechanically checked against generated trace artifacts.

1. Introduction

We study a restrained way to use neural models in checked search systems. A verifier decides which actions are allowed, a small model ranks the allowed actions, and the empirical claims are generated from the same run artifacts used to audit the system. The contribution is not that tiny networks solve hard tasks. It is that checked failures and checked successes become reusable training and proposal structure without giving the model authority over correctness. The clean anchor is arithmetic sequence filtering: in the audited next-prime instance (`prime_gap` in traces), the generated claim row reports the verified learned filter raises top-1 ranking to 68.87% from 45.12% for the hand-coded wheel and 9.13% raw, with 0 verifier-counted true-answer deletions. The same family also admits recurrence filters such as Pisano-period residue gates, but the strictly audited portion remains the next-prime gap under segmented-sieve verification.

The dataset-first engine accepts new experiments as candidate rows plus verifier labels, generating evidence artifacts before any prose claim is written. Two architectural extensions are under development and described in companion reports: holographic weight storage [1] for graceful degradation and content-addressable weight lookup, for improved neural ranking on structured continuous features. This extension is compatible with the framework but not required for the results presented here.

The rest of the paper is organized around the data flow. Section 2 positions the work against related approaches. Sections 3-4 define the verifier-grounded protocol and dataset engine. Section 5 presents the frozen evidence matrix. Section 6 provides the arithmetic filtering anchor. Section 7 covers the holographic memory substrate. Section 8 Sections 8-9 cover feedback loops and boundary evidence. Sections 11-12 cover claim governance and limitations.

2. Positioning

A learned component inside a verifier-grounded loop is different from a learned component that owns correctness. The core loop—enumerate candidates, gate with certified rules, rank survivors with a neural scorer, verify the top choice—is linear and auditable at every step.

The gated ranking loop has three comparison frames. **Verification-first systems** [2] rely on sound checkers but few close the loop on proposal structure. **Neuro-symbolic scoring** [3] uses LLM ranking without the scoring-gating feedback that turns negative rows into training signal. **Code-generation with solvers** [4] uses neural proposal around formal environments. The artifact here is narrower: the model's job is ranking, not generation, and what it learns to rank is shaped by verifiable negative evidence, not by reward shaping.

The holographic memory substrate draws on vector symbolic architectures [5] and holographic reduced representations [6]. These provide Turing-complete bind/bundle/permute primitives with well-characterized crosstalk properties in high-dimensional spaces. The leOS-C engine [7] isolates these primitives as a compact C library that can be linked directly into the scorer pipeline.

3. Protocol

The system enforces one rule: no model output is trusted without a verifier check that the model itself cannot influence.

3.1. The Certificate Gated Loop

Every decision follows five steps. (1) **Enumerate**: given a state, generate the full candidate action set. (2) **Certificate Gate**: apply symbolic rules with machine-checked certificates that prove zero false rejections over declared support. (3) **Rank**: a neural scorer assigns a score to each survivor. Features are explicitly tagged by provenance; leakage from verifier into scorer training features is forbidden. (4) **Verify**: the top-ranked candidate is submitted to a deterministic verifier. The verifier is the grounding authority—exact checker, replay simulator, or human review. (5) **Trace**: every step is recorded with stable action IDs. The trace becomes the evidence artifact for claims.

3.2. Feature Envelope and Leakage Policy

Every candidate action is encoded as a 96-dimensional feature row. Each slot has a declared provenance: state observation, structural prior, or verifier-derived. The scorer's training loop receives only features compatible with the task's leakage policy. Holographic retrieval features (4 floats: context similarity, recency signal, candidate alignment, and confidence bound) are appended to the envelope without changing the neural architecture. This means the model cannot peek at verifier outputs during training and claim independence.

3.3. Rule Grammar and Certificates

Symbolic rules are governed by a domain-specific grammar. Each rule has a kind (reject_action, transform, threshold), typed parameters, and a certificate. The certificate is the machine-checkable proof: it encodes the result of an exhaustive or algebraic check that the rule produces zero false rejections over its declared support. Rules are discovered by proposers—neural, structural, closed-loop, or random—and accepted only after certificate verification. Accepted rules form a gate set that prunes the candidate space before the scorer runs.

4. Dataset-First Engine

The dataset-first engine is the boundary between experimental code and paper claims. Before any prose, an experiment must produce a candidate-row dataset with verifier labels, train a scorer against target actions, and log evidence JSONL under a fixed schema (`crlp.dataset_action_score.summary.v1`).

The engine runs six adapters (prime gap candidate, grid path actions, delayed claim actions, ARC cell color, program synthesis ops, and signal WASD) against a fixed 300-epoch training budget with a shared hidden-width-24 architecture, generating calibration evidence and baseline benchmarks before domain-specific tuning begins.

This constraint is a quality gate, not a restriction. Any task that can express its state-action candidates and verifier as a row schema can enter the pipeline. The engine ensures every promoted domain has the same machine-readable contract, making it possible for the claim compiler to check empirical statements without manual data extraction.

5. Evidence Matrix

The promoted spine carries machine-checked claims. Sandbox domains carry diagnostic claims that track progress toward promotion.

prime_gap. Certified modular filters shrink the residual action space. (Verifier: segmented-sieve replay and modular certificate checks.)

grid_path. The same gate contract works outside residue arithmetic, applying legality, topology, and distance-dominance certificates to spatial pathfinding. (Verifier: BFS shortest-path successor and rule certificate checks.)

delayed_claim. Reason learning needs domain-aware safety constraints: threshold and margin predicates operate over observable claim features under a deterministic hidden-outcome simulator. (Verifier: deterministic replay and finite-generator safety checks.)

rejected_candidate_learning. Verifier failures become training signal. (Verifier: original rule checker that rejected or accepted each row.)

hypothesis_library_feedback. Certified learning accumulates through reusable library priors, with zero-false-rejection promotion gates. (Verifier: train-pair exactness, type checks, and held-out fixture evaluation.)

arc_agi, program_synthesis, haiku_word_lm, signal_flight. These sandbox domains publish the same schema contract but carry diagnostic claims. They demonstrate generalization beyond arithmetic domains while remaining outside the canonical evidence surface until earning broader verifier-backed coverage.

6. Anchor: Arithmetic Sequence Filtering

The cleanest test case is next-prime gap prediction. Candidate actions are even gaps 2..256 plus an overflow action. Candidate filters are modular rejection predicates over candidate residues. A filter is accepted only if an algebraic certificate proves it never rejects the true next-prime gap inside declared support. After gating, the scorer ranks remaining candidates.

Condition	Top-1	Top-10	False rejections
no gate	9.13%	78.01%	0
hand wheel	45.12%	99.77%	0
learned wheel	**68.87%**	**100%**	**0**

Table 1: Arithmetic sequence filtering on 30000 rows: the verified learned filter raises top-1 ranking to 68.87% from 45.12% for the hand-coded wheel and 9.13% raw, with 0 verifier-counted true-answer deletions. Learned-vs-hand top-1 delta: +23.74 pp.

The seed audit checks that the headline does not depend on one lucky training run: across 5 seeds, the learned filter stays at 68.87% +/- 0.00pp, versus 45.12% +/- 0.00pp for the hand wheel and 9.13% +/- 0.00pp raw; the worst learned-vs-hand gain is +23.74pp, learned top-10 never falls below 100.00%, and false rejections never exceed 0; promotion requires 5 seeds. The zero spread is a seed-stability audit over this deterministic benchmark: seeds vary checkpoints and proposer runs, not the declared prime-range evaluation distribution.

The learned proposer adds a useful negative row. The naive learned proposer fails: the static fixed-budget proposer is a supported negative baseline: candidate static saves 72.49 vs best deterministic 73.22, and conditional static saves 2.15 vs best deterministic 55.57. That failure prevents the claim set from becoming a generic "neural proposal wins" story. The naive learned proposer fails—that is the point. Unsupported neural proposal is not trusted; verifier feedback is recycled until proposal economics improve. In this protocol, negative rows are not cleanup work after the fact. They are the training substrate that separates model ranking from validity.

7. Ongoing Work

Two architectural extensions are under active development and described in companion technical reports.

Holographic weight storage [1] replaces flat weight arrays with vector-symbolic bindings on the unit hypersphere, giving neural scorers graceful degradation (1/sqrt(N)), content-addressable weight lookup, and fixed O(d) memory. A zero-parameter holographic ranker achieves 48.5% accuracy versus 16.0% for a gradient-trained MLP on small-sample tasks.

universal primitive for continuous functions, analogous to NAND for digital logic. Integrated into neural scorers as pairwise feature enrichment, it improves accuracy pairwise layers compose correctly at 94.5%. The operator does not yet improve the promoted experiments, which are dominated by binary and categorical features.

This extension is compatible with the certified-gate framework. Neither is required for the results presented in Sections 2–10.

8. Feedback Loops

The main contribution is what happens after the checker has said yes or no. Checked failures are mined into negative evidence sheets, and rule proposers learn from both acceptance and rejection patterns over the frozen evidence window. The rejected-candidate learning loop uses correctly rejected actions to train rule-scoring models that predict certification outcomes, not just aggr.certificate_gate ranking scores. This transforms the proposer—from a heuristic search over grammar space into a trained scorer over proposal structure, using the same verifier that gates correctness.

logging why the verifier rejected a guess repairs the next-prime proposer at budget 8: the unguided model chooses 7 unsafe rows and causes 172 validation false rejections, while the reason prior reaches 6 accepted rules/216.0 utility and the trained repair reaches 7 accepted rules/245.0 utility with zero false rejections; arithmetic-sequence/modular split sweeps pass, but the delayed-claim stress case has 0 mixed split rows. The key metric is that the system converts check—and-reject cycles into training signal for the next proposal round. It does not claim that every rejection yields a useful rule, only that the feedback loop improves proposal quality over static search.

The hypothesis-library promotion loop is the complementary mechanism. Rules that pass certification are not stored in an inert log. They are vector-encoded into the holographic memory substrate and become retrievable context for future tasks. A task signature query against the holographic store returns

a weighted blend of all previously learned rules relevant to the current task in one vector look-up, enabling cross-task transfer without explicit rule sharing. saving repeatedly safe rules helps future tasks: promoted move_to_anchor from 2 tasks/78 covered events with zero false rejections; slot priors raise first-choice certifications from 2 to 4 and budget-8 utility from 0.2119 to 0.5169. Held-out split feedback rows: 30.

9. Boundary Evidence

The evidence matrix is designed to prevent overclaiming. Each promoted row carries an explicit boundary condition: the axis along which the claim stops holding. Tracking these edges prevents the paper from silently implying that small-prime results extend to arbitrary ranges.

Tiny-net plateau. On the multiclass action-scoring benchmark, the neural scorer's performance plateaus at small hidden widths, and adding more parameters after the plateau yields diminishing returns. tiny-net plateau audit covers 6 capacities, 3 restarts, and 8 checkpoints through 128 epochs across 12 domain/variant verdicts; 20/72 required rows meet the plateau threshold and 0/12 verdicts are audit-ready.

Trace determinism. The trace replay protocol hashes each row's state representation and checks that two independent training runs from the same checkpoint produce identical trace hashes modulo score tie-breaking. the trace audit checks whether replay meaning survives numeric/runtime changes: 2 traces and 2000 rows, 0 invalid rows, 0 rank-sensitive rows, 0 logged mismatches, same semantic hash: no.

Dataset engine evidence surface. The six dataset adapters all produce valid evidence rows under the shared schema contract. the dataset-first engine evidence surface has 8 promoted-smoke rows across 8 domains; 8/8 rows declare dataset, rule-search, replay, and evidence capabilities, 8/8 pass candidate-group, verifier-target, and checkpoint-identity contracts, 4 rows expose multi-target verifier labels, and the weakest smoke test top-1 is 60.00%. These rows are infrastructure evidence, not promoted domain-result claims.

Signal flight. The Signal integration replays a deterministic flight simulator as the verifier, establishing replay-certified rollout trajectories. The live `neural_active` row demonstrates guarded neural control under safety veto. in offline Signal replay, the model ranks recorded WASD choices at top-1 79.76% vs baseline 67.86% (+11.90pp), utility 0.1432 vs 0.1359, regret 0.0073 vs 0.0146, damage 2.38% vs 2.38%, deaths 0.00% vs 0.00%; live shadow frames=74. The row establishes live end-to-end execution with safety gating but does not show comparative neural improvement over a heuristic baseline. the live Signal pilot did run checkpoint-ranked neural actions, but only behind guards: 74 recent neural-source rows from checkpoint f12b644ed9, p95 inference 0.089 ms, 74 replay-ready rows, protocol-compatible rows=74, safety-gate rows=74, damage events=2, deaths=0; no baseline comparison is claimed.

10. Claim Governance

Claims are not prose. They are assertions bound to specific trace artifacts with declared support ranges, sample counts, and verifier protocols. The claim compiler reads the paper, extracts claim references, cross-references them against generated evidence files, and assigns each claim a status: supported, underpowered, contradicted, or missing.

```
(
  supported: 14,
  contradicted: 1,
  underpowered: 5,
  missing: 0,
```

```
  stale: 0,
)
```

Claims that are supported appear in the results sections above. Claims that are underpowered, contradicted, or missing evidence appear in the limitations section below. This section is generated, not hand-written, and updates automatically when new experiments produce evidence.

C7 blocks this capacity claim until plateau audit is ready: the sweep's best fixed-ARC shared row is 32x16 at 39.20% and delayed_claim deltas are +0.10pp to +0.43pp, but only 20/72 plateau rows and 0/12 domain/variant verdicts are audit-ready. Adapter heads show that task-specific fine-tuning recovers performance after shared-trunk multitask training depresses it, but the claim is restricted to diagnostic capacity measurement rather than a promoted multitask architecture.

11. Limitations

```
(
  (
    claim_id: "C4.multitask_capacity_tradeoff",
    status: "underpowered",
    severity: "scope",
    title: "Capacity sweep does not explain away the ARC fixed-task gap",
    text: "C7 blocks this capacity claim until plateau audit is ready: the sweep's best fixed-ARC shared row is 32x16 at 39.20% and delayed_claim deltas are +0.10pp to +0.43pp, but only 20/72 plateau rows and 0/12 domain/variant verdicts are audit-ready.",
    roadmap_items: (),
    residual: (),
  ),
  (
    claim_id: "C4.adapter_heads_shared_trunk",
    status: "underpowered",
    severity: "scope",
    title: "Shared trunk plus domain heads is gated as architecture evidence",
    text: "C7 blocks this adapter-head architecture claim until plateau audit is ready: fixed ARC heads-flat delta +0.00pp; delayed_claim heads-flat +0.00pp; grid_path heads-flat +0.23pp; ARC OOD heads-single +5.80pp top-1 and -1.2020 NLL, with only 20/72 plateau rows and 0/12 domain/variant verdicts audit-ready.",
    roadmap_items: (),
    residual: (),
  ),
  (
    claim_id: "C7.tiny_net_plateau_stop_rule",
    status: "underpowered",
    severity: "scope",
    title: "Tiny-net architecture claims are gated by plateau audits",
    text: "tiny-net plateau audit covers 6 capacities, 3 restarts, and 8 checkpoints through 128 epochs across 12 domain/variant verdicts; 20/72 required rows meet the plateau threshold and 0/12 verdicts are audit-ready.",
    roadmap_items: (),
    residual: (),
  ),
  (
    claim_id: "C8.trace_determinism_audit",
    status: "contradicted",
    severity: "scope",
```

```
    title: "Trace determinism audit catches replay-sensitive claim drift",
   text: "the trace audit checks whether replay meaning survives numeric/runtime changes:
2 traces and 2000 rows, 0 invalid rows, 0 rank-sensitive rows, 0 logged mismatches, same
semantic hash: no.",
    roadmap_items: (),
    residual: (),
  ),
  (
    claim_id: "C5.arc_verified_self_play",
    status: "underpowered",
    severity: "scope",
    title: "ARC verified self-play is a fixed-budget search-efficiency claim",
      text: "ARC verified self-play is registered but underpowered: missing build/
float/arc_verified_self_play/summary.jsonl, so no public replay, fixed-budget ablation,
leakage audit, or zero-family opening can be credited.",
    roadmap_items: (),
    residual: (),
  ),
  (
    claim_id: "C9.signal_mining_replay_options",
    status: "underpowered",
    severity: "scope",
     title: "Signal mining option-ranker evaluates bounded policy hypotheses through
replay",
      text: "Signal mining claim registered as underpowered pending summary artifact
generation.",
    roadmap_items: (),
    residual: (),
  ),
)
```

The promoted spine covers three task families with exact-check and replay verifiers. The sandbox tasks (ARC, program synthesis, haiku, signal) are experimental proving grounds. Extending coverage requires broader verifiers: visual-algebra programs for ARC, DSL correctness checkers for program synthesis, and semantic understanding transfers for haiku.

The holographic memory substrate is demonstrated at d up to 2048 with hundreds of stored rules. Scaling to millions of rules requires dimension scaling (crosstalk scales as 1/sqrt(d)) or hierarchical VSA structures. Holographic memory recall requires linear superposition; the bind/unbind primitives are well-characterized in this regime, and alternative binding operations should be confined to representation-learning layers rather than the memory storage path.

The dataset-first engine currently enforces a fixed 300-epoch, hidden-width-24 baseline. While this provides a consistent calibration point, it does not represent the optimal configuration for any individual task.

12. Conclusion

We presented a verifier-grounded architecture for neural search. The core loop is simple: enumerate candidates, gate with certified rules, rank survivors with a neural scorer, verify the top choice, and preserve the trace. The claim compiler closes the loop between empirical work and scientific reporting by mechanically checking every assertion against generated trace artifacts.

The architecture is honest about its limits. Certified gates do not claim to operate outside their declared verifier support. The learned proposer's failure (the static fixed-budget proposer is a supported negative baseline: candidate static saves 72.49 vs best deterministic 73.22, and conditional static saves 2.15 vs best deterministic 55.57.) is documented as a negative row. The dataset-first engine ensures every promoted domain publishes the same machine-readable contract. The claim compiler ensures the paper never says more than the evidence supports.

The dataset-first engine makes the pipeline repeatable: any experiment that exports candidate rows and verifier labels under the shared schema can enter the evidence matrix. Two extensions under development—holographic weight storage framework but not required for the results presented here.

Bibliography

[1] Anonymous, "Holographic Weight Storage for Verifier-Grounded Neural Search." 2026.

[2] R. Alur and others, "Automata-Based Verification for Neural Systems," 2024.

[3] L. Gao and others, "Program Synthesis with Large Language Models," 2024.

[4] S. Polu and I. Sutskever, "Generative Language Modeling for Automated Theorem Proving," *arXiv preprint arXiv:2009.03393*, 2020.

[5] D. Kleyko *et al.*, "Vector Symbolic Architectures as a Computing Framework for Nanoscale Hardware and Emerging Applications," *Proceedings of the IEEE*, vol. 110, no. 10, pp. 1538–1571, 2022, doi: 10.1109/JPROC.2022.3209104.

[6] T. A. Plate, *Holographic Reduced Representations*. CSLI Publications, 2003.

[7] rati, "leOS-C – Holographic Engine in C." [Online]. Available: https://github.com/ratimics/leos-c